

Summarizing Audit Trails in the Aeouls Security Platform

by

Wissam Jarjoui

S.B., C.S M.I.T., 2011

Submitted to the Department of Electrical Engineering and
Computer Science

in partial fulfillment of the requirements for the degree of
Master of Engineering in Electrical Engineering and Computer
Science

at the

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

September 2012

© Massachusetts Institute of Technology 2012. All rights reserved.

Author.....
Department of Electrical Engineering and Computer Science
August 15, 2012

Certified by
Barbara H. Liskov
Institute Professor
Thesis Supervisor

Accepted by
Christopher J. Terman
Chairman, Masters of Engineering Thesis Committee

Summarizing Audit Trails in the Aeolus Security Platform

by

Wissam Jarjoui

Submitted to the Department of Electrical Engineering and Computer Science
on August 15, 2012, in partial fulfillment of the
requirements for the degree of
Master of Engineering in Electrical Engineering and Computer Science

Abstract

For an Aeolus user, analyzing an Aeolus event log can prove to be a daunting task, especially when this log grows to include millions of records. Similarly, storing such an event log can be very costly. The system I present in this thesis provides an interface for the creation of user-defined summaries of the Aeolus audit trails, as well as truncation of the event log, making it easier for to analyze the Aeolus event log and less costly to analyze events of interest. This is done through through the use of what we term a *Summary System* and *Summary Objects*. I present the system in the context of a sample application based on *www.mint.com*. The system is an extension to the Aeolus library, it is written in Java and uses a PostgreSQL database as it's primary database.

Thesis Supervisor: Barbara H. Liskov
Title: Institute Professor

Acknowledgments

First and foremost, I would like to thank my advisor, Prof. Barbara Liskov, for the mentorship she gave me in my two years here at PMG. Her guidance helped me navigate and learn a lot about systems security, and went beyond just this project. I am truly grateful for the understanding, patience and investment she put in while helping me see this project through; I hope that one day I would be able to reflect those traits as a leader.

A big thank you to my colleague, David Schultz, for his tremendous help in this project and my personal learning; from suggesting and brainstorming ideas to improve my project, to explaining technical concepts, and being a strong source of knowledge and inspiration and a good role model in general. David, it has been a pleasure working with you.

I would also like to thank James Cowling, Dan Ports and Barzan Mozafari and everyone in 32-G908 for all the fun conversations and useful advice they provided. I am very fortunate to have had the chance to hear and learn from their experiences, and be inspired by them.

As this thesis culminates five years for me at MIT, I would like to take this chance thank my friends and family. My friends, for being there for me, and for offering me help even when I didn't ask. Their presence was a home away from home, and their support guided me in times of uncertainty. My family, without whom I would not be where I am today, for their support and encouragement, for worrying about the things I wouldn't, and for setting the bar high. In one way or another, I will always have something to learn from them.

Finally, I would like to thank MIT for all what it has given me, in both computer science and real life. I thank Anne Hunter, my advisors, my professors and my TAs for holding my hand when I took on different challenges. I thank the Stata Center, the Green Building, and The Infinite Corridor for being the monuments of my experience here at MIT, and I thank the Charles River, for being the source of calmness in a very busy time. They will all be missed.

Contents

1	Introduction	6
1.1	Motivation for Summarization and Archiving	7
1.2	Thesis Outline	8
2	Aeolus System Overview	9
2.1	System Architecture	9
2.2	Information Flow Model	10
2.2.1	Prinicpals, Tags and Labels	10
2.2.2	Information Flow Rules	10
2.2.3	Authority	11
2.2.4	Compound Tags	12
2.3	Programming Model	12
2.3.1	Threads and Vitual Nodes	12
2.3.2	Shared State Objects and Boxes	13
2.3.3	Authority Closures and Reduced Authority Calls	13
2.3.4	Files	13
2.3.5	Log Collection	14
3	Mint	16
3.1	Mint Model	16
3.1.1	Authority State Model	17
3.1.2	Files	19
3.1.3	Shared Memory Objects	21

3.2	Implementation	22
3.3	System Security	24
4	Auditing the Mint Application	26
4.1	Detecting Suspicious User Activity	29
4.2	Summarization	31
5	The Summary System	32
5.1	Summarization Model	33
5.2	Summarization Workflow	34
5.3	The QUERY Table	34
5.3.1	The Query Management Interface	35
5.3.2	Query Attributes	36
5.4	Producing Summaries and Marking Events	36
5.4.1	SummaryObjects	37
5.4.2	Summary Attributes	38
5.5	Marking Events	39
5.6	Discussion	39
6	Auditing Using Summarization	41
6.1	Summarizing User Sessions	41
6.2	Summarizing File System Events	45
6.3	Summarizing User Trends	46
6.4	Sufficiently Summarized Information	48

Chapter 1

Introduction

Security of confidential online information, such as medical records and financial data, is a very important problem. Recent research has focused on Decentralized Information Flow Control (DIFC) as the most promising approach to enable application developers to secure information. DIFC is based on the principle that the system tracks information as it flows through the system, and only allows release if the releaser has sufficient authority. Additionally DIFC provides this ability in a fine-grained way, so that security policies can be tailored to the needs of individual users and organizations.

This thesis extends the security support provided by the Aeolus platform. Aeolus is a platform that combines DIFC and an intuitive security model framework to make it more convenient for developers to build secure applications on a distributed system. Additionally, Aeolus provides automatic auditing of every security related event that occurs while an application runs; furthermore, it provides a way for applications to log additional events that are meaningful at the application level. The audit trail is an important component of security to allow discovery of errors that cause security policies to be subverted; the audit trail can also be used to discover attacks.

An issue in any auditing system, including the Aeolus system, is that the audit trails can become extremely large. This is especially true if the system being audited is large, e.g., has millions of users, and long-lived. Therefore a way of re-

ducing the stored information and making the important information more easily accessible is needed.

This thesis addresses this problem. It provides a framework that allows groups of audit events to be summarized: the important content of the group of events is captured in a *summary event*, which is much smaller than the group. Once information has been summarized, the base events that underlie the summary can be moved to archival storage, or even deleted. An additional benefit is that summary events are more accessible than the base events they summarize because they can be defined to explain the information in application-specific terms.

A final point is that the production of the summary events is itself controlled by information flow. This is important because, everything a system does is detailed in the audit log, and therefore, there is a potential for information leaks if the log could be accessed by an unauthorized party.

The next two sections describe the motivation and scope of this system with the help of some real-life examples.

1.1 Motivation for Summarization and Archiving

Summarization allows developers to group information that is spread out over the Aeolus audit trails into smaller space. This allows for future archiving or truncation of the Aeolus log, as well as setting the stage for writing more faster queries.

Take for example a bank's web administrator who is interested in detecting suspicious activity on customers' bank accounts. One way to do this is to produce a plot of the number of outbound transfers carried through a user's account for each week in the last year. Any spikes in the rate of outbound transactions per week could mean that a user's account has been compromised by an attacker. Finally, the administrator might also wish to produce such a plot at the end of each week.

Producing such information is possible through the current auditing mechanisms available in Aeolus. However, in order to accomplish such a task, the administrator's audit will have to scan all events in the system in the past year, which

could span millions of events, once a week, and this would be costly.

Summarization solves this problem by allowing the web administrator to store weekly computations as summaries themselves, and simply reusing these summaries in future weekly audits. For example, the web administrator could store the average number of outbound transfers a user has made per week in the last year and use that as a benchmark for future user activity ¹.

1.2 Thesis Outline

The remainder of this thesis is organized as follows: Chapter 2 presents an overview of the Aeolus security system. Chapter 3 presents a sample Aeolus application based on the financial management service Mint.com. Chapter 4 describes the summarization models and interfaces. Chapter 5 describes the implementation details of the system. Chapter 6 evaluates the ease of use of the system as well as overall system performance. Chapter 7 discusses related work in streaming databases. Chapter 8 presents some topics for future work and reviews the contributions of this thesis. Appendix A details the summarization and archiving interface and presents some examples.

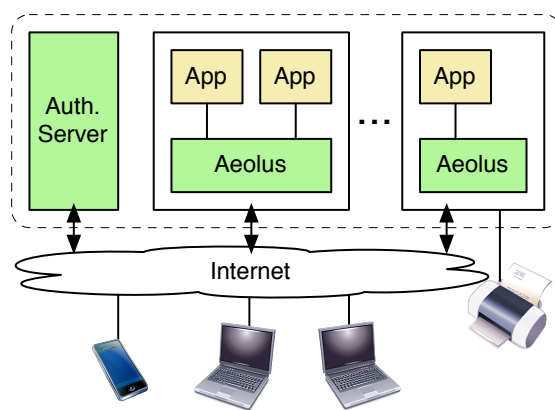
¹These summaries could be considered *summaries of summaries* as they further summarize previously-summarized information.

Chapter 2

Aeolus System Overview

This chapter presents an overview of the Aeolus security platform, on which our summarization system is based, and highlights relevant details such as log collection. More complete descriptions of the Aeolus security platform and its log collection and analysis can be found in Cheng [3], Popic [5] and Blankstein [1].

2.1 System Architecture



Aeolus System Architecture

Figure 2-1: High level overview of Aeolus system architecture.

The Aeolus architecture is shown in figure 2-1. The system consists of many nodes, each of which is trusted to enforce Aeolus's information flow rules.

Aeolus tracks information flow within each system node and between system nodes. Nodes in the system communicate via RPC messages, those messages are encrypted to protect their secrecy and integrity. Nodes outside the system are considered to be untrusted: information can flow outside of the system only if it is uncontaminated, and information arriving from the outside is marked as having no integrity.

Threads in Aeolus run on behalf of *principals*. The ability of a thread to carry out *privileged operations* is determined by the *authority* of the principal it runs on behalf of. Aeolus tracks the authority of principals in the *authority state* stored at the authority server (AS). Principals, privileged operations and authority are described in sections 2.2.1, 2.2.2 and 2.2.3, respectively.

2.2 Information Flow Model

This section describes the basic concepts and rules of the Aeolus security model.

2.2.1 Principals, Tags and Labels

Aeolus employs an intuitive security model to implement information flow control. The model revolves around three key concepts: *principals*, *tags* and *labels* [2]. Principals represent entities in the system that create, modify and share information. Tags represent security categories of information. A principal authoritative for a certain tag can modify or share information categorized by that tag. Labels are sets of tags and are used to determine whether information can flow from a source to a receiver based on the information flow rules described in the following section.

2.2.2 Information Flow Rules

Aeolus allows information to flow from a source S to a destination D only if the following rules are satisfied:

$$\begin{aligned} \text{SECURITY}_S &\subseteq \text{SECURITY}_D \\ \text{INTEGRITY}_S &\supseteq \text{INTEGRITY}_D \end{aligned}$$

Threads in an Aeolus node run with security and integrity labels associated with them; objects such as files also have such labels. In Aeolus, a thread cannot modify labels of files or shared memory objects; however, the labels of threads can change, and hence a thread would have to modify its own labels in order to read or write data.

Certain label manipulations, called *privileged manipulations* are unsafe because they remove constraints on information flow:

1. *Declassification* Remove a tag from a secrecy label.
2. *Endorsement* Add a tag to an integrity label.

A thread is allowed to carry out privileged label manipulations if it is running on behalf of a principal that has authority for the tags affected by these manipulations.

2.2.3 Authority

Authority determines whether a thread can perform privileged label manipulations. Authority starts with tag creation: when a thread creates a tag, its principal has authority for that tag.

Subsequently, the Aeolus authority state can be modified through *grant*, *act-for* or *revoke* operations. Grant operations allow a principal to delegate authority for a particular tag to another principal. *Act-for* operations allow a principal to delegate all of its authority to another principal. Revoke operations remove *act-for* and *grant* links between principals. To avoid covert channels, Aeolus only permits threads running with null secrecy labels to modify the authority state.

2.2.4 Compound Tags

Applications frequently have sets of tags that are closely related. In order to simplify the authority structure of the application, Aeolus allows for tags to be grouped upon creation using *compound tags*. For example, in a medical clinic, patient data tags are *subtags* of an *ALL-PATIENT-DATA* tag, a *supertag*.

A principal authoritative for a supertag is also authoritative for all of its subtags. Similarly, having a supertag in a thread's secrecy label is equivalent to having all subtags in its secrecy label. This reduces label size substantially, and makes label manipulations involving particular groups of tags less expensive.

See figure 3-1 on page 17 for an example of an authority state model, detailing *act-for*, *grant* and *subtag* relationships in the financial services application described in chapter 3.

2.3 Programming Model

This section explains the programming abstraction Aeolus provides, and how they support DIFC.

2.3.1 Threads and Virtual Nodes

Virtual nodes (VNs), shown as applications in figure 2-1, communicate via RPC and have many threads inside them. A VN is given a principal when it is created and all threads run with this principal or some principal it acts for. An RPC is run in its own threads with the VN principal but the labels of the caller; on the return the labels are sent back to the caller where they are merged with those of the caller: a merge is a union of the secrecy labels and an intersection of the integrity labels. Then the caller continues with its own principal.

2.3.2 Shared State Objects and Boxes

Within a virtual node, threads can share state via special Aeolus shared objects. Aeolus gives threads access to a *root* shared memory object, with null labels.

Aeolus shared objects have immutable labels, and Aeolus ensures that the rules in 2.2.2 are respected. Users can create their own shared objects; Aeolus ensures that label manipulations do not take place inside the object's methods.

2.3.3 Authority Closures and Reduced Authority Calls

Aeolus provides developers with *authority closures*. An authority closure is an object bound to a principal at the moment of creation. Threads can later call the closure's methods, which run with the closure's principal. Authority closures allow threads to process confidential information without being exposed to the information itself.

Aeolus also provides a mechanism for threads to reduce their authority, called *reduced authority calls*. To make a reduced authority call, a thread specifies a function and a principal to run it with. The calling thread's principal must act for the principal of the reduced authority call.

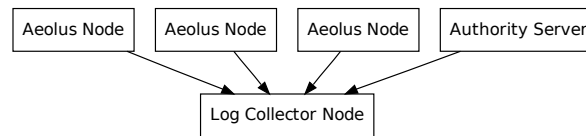
With those two mechanisms in place, Aeolus makes it possible to ensure that each part of the application runs with only the authority it needs. Aeolus also provides a principal that is not authoritative for any tags, P_{PUBLIC} , for developers to use when they need to ensure that a part of a program cannot leak any information.

2.3.4 Files

Aeolus provides a network file system that enforces DIFC. Similarly to shared memory objects, files have immutable labels and the rules in 2.2.2 apply for information flow in and out of them.

Files are yet another way (RPCs, shared memory objects) through which Aeolus allows threads to communicate. A complete description of the Aeolus file system API is presented in McKee [4].

2.3.5 Log Collection



Log Collection

Figure 2-2: Flow of logs to the log collection system.

Aeolus provides automatic auditing of every security related event that occurs while an application runs; also it provides a way for applications to log additional events that are meaningful at the application level.

Aeolus distributes log collection across all nodes in the system. Aeolus sends events that occur locally to the log collector as shown in figure 2-2. The log collector stores the events for later processing and analysis, as explained in [1].

Aeolus provides application-level logging to allow developers to log their own events using the following interface:

```
AeolusLib.createEvent(String app_op, List<String> args)
```

This creates a new record in the Aeolus event log. The *app_op* and *args* argument are stored as attributes for that record, allowing the developer to query the event log based on the arguments they provided. The *app_op* parameter is used to identify the application operation name, and the *args* parameter allows the developer to specify any additional information to store for the event.

Aeolus Event Attributes

Aeolus stores different attributes for different events that take place in the system. Blankstein [1] provides a complete description of those attributes; here we will highlight the ones more relevant to the use of summarization:

event_counter

The event counter uniquely identifies an event, and provides an ordering for

the occurrence of the event, i.e. events with a higher eventcounter took place after events with a lower eventcounter.

timestamp

This is the real time at which the event happened.

secrecy

This is the secrecy label of the thread which caused the event, just before the event took place.

integrity

This is the integrity label of the thread which caused the event, just before the event took place.

op_name

This field specifies the type of the event, e.g. a *DECLASSIFY* event, or a *SENDERPC* event. Those types are detailed in [1]. Application-level events created with the above events all have the *APPLICATION_LVL* operation name.

app_op

This field holds the name of the application-specific operation and is provided by the developer as described above.

app_args

This field holds the arguments specified for an application-level event (the second argument in the interface shown above).

running_principal

This field holds the principal of the thread that caused the event.

It is important to note that the *secrecy* and *integrity* attributes are particularly important to prevent information leakage in the system: a thread can only read those event records if the DIFC rules in 2.2.2 are satisfied.

[mention this ideally runs on david's db]

Chapter 3

Mint

Prior to the development of the summarization system, it was important to develop a sense for what is needed from such a system, and what problems can it solve. To answer these questions, I developed an Aeolus application based on the financial management service *mint.com*. The *Mint* application served as a case study and guided the development of the summarization system.

In this chapter I present the Mint application model as well as some examples to how a summarization system could be used for this application.

3.1 Mint Model

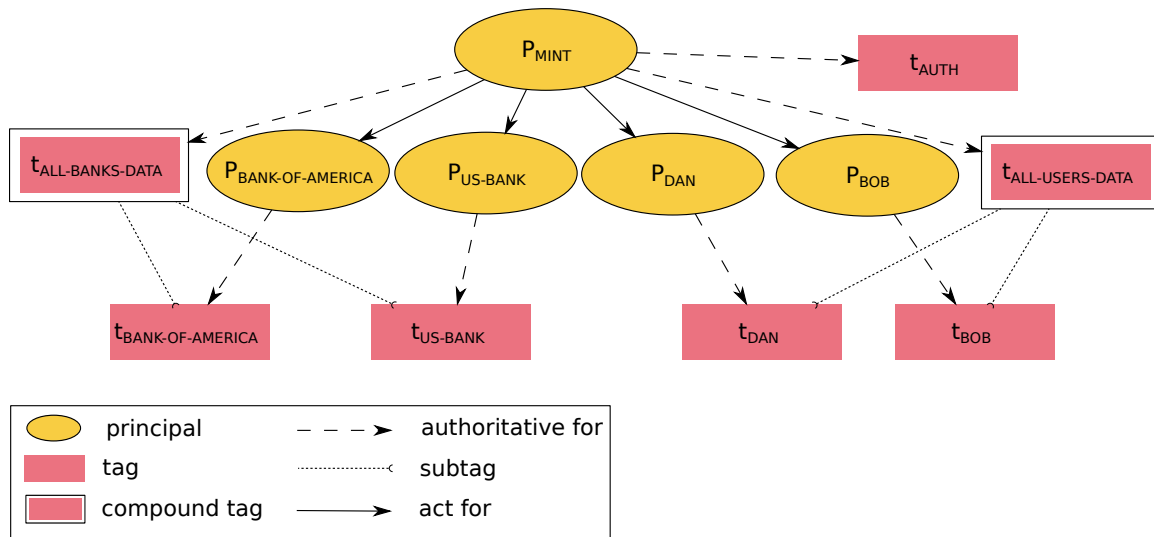
Mint.com is an financial management service that provides its users with the ability to monitor their bank accounts across different banks from one place. It also allows users to run transaction analysis tools that will access the information of certain transactions for each bank for the user, and generate a result. In addition, Mint runs aggregate analysis for all user accounts.

For example, Bob could sign up on www.mint.com and add his Bank of America account and his US Bank account for monitoring. This will cause our application to store credentials for both of these banks for Bob. Bob can then request a graph that presents his expenditures on food in the last week. Bob's actions would also affect the general aggregate analysis done over all Mint users.

This example gives us a starting idea of how should we build our authority structure. For instance, there should be a *BOB-DATA* tag associated with Bob's information (e.g. expenditure statistics, transactions, etc.) that only a principal running on Bob's behalf can declassify. Note that Bob's password should never be revealed to Bob himself, neither should any of Bob's bank credential be revealed as well. Furthermore, there's a natural grouping of user data tags, and hence a super tag, *ALL-USER-DATA*, could be useful to run the aggregate analysis tools for Mint.

In this section we examine the security model of the application more closely.

3.1.1 Authority State Model



Mint Authority State Model

Figure 3-1: This diagram illustrates the authority state model of the Mint application in an example where two banks are registered at the application, Bank of America and US Bank. Users Bob and Dan have both signed up. Principals are denoted $P_{\langle name \rangle}$ and tags are denoted $t_{\langle name \rangle}$.

In this section we identify the principals running in the system, the tags associated with different data, and the relationships between them. Figure 3-1 shows an overview of the authority state model for Mint.

An intuitive way to list all necessary¹ principals, is to think of the different

¹Necessary by good design. The application could run with just one principal.

“clearance levels” in the system: some information should be accessible by a user, some by a bank, and some information should not be accessible at all. In the light of these requirements, the application employs the following principals:

Mint principal

This principal is authoritative for all tags in the system, and is used for aggregate analysis of user data and user authentication.

Bank principals

Each of these principals is authoritative for a particular bank’s tag.

User principals

Each of these principals is authoritative for a particular user’s tag.

Public principal

Described in section 2.3.3, this principal is used when no privileged operations are necessary to ensure no information leaks.

Similarly, to separate data in the system according to its purpose, (e.g. user password, credentials to download bank transactions, etc.), the application employs the following tags:

User data tags

Each of these tags is associated with data carrying a particular user’s information. Reading (or modifying) user information requires declassifying (or endorsing) a user data tag.

Bank data tags

Each of these tags is associated with data carrying a particular bank’s information. Reading a user’s bank credentials requires declassifying a bank data tag. This tag is important so that only a Bank’s closure is able to access a user’s bank password.

All-Users-Data tag

This tag is a super tag for all the user tags. Running aggregate analysis on

Mint user accounts requires declassifying this tag.

All-Banks-Data tag

This tag is a super tag for all the bank tags. Adding a new bank to Mint requires endorsing this tag.

Authentication tag

This tag is associated with user's Mint passwords. Authenticating a user's mint credentials requires declassifying this tag. This tag is important to prevent release of a user's password to someone who might have gained access to their account.

The following two sections explain how data is stored in the system and which tag operations are carried out in a typical workflow.

3.1.2 Files

Files	Secrecy Label	Integrity Label
mint-dir	{}	{ $t_{\text{ALL-USERS-DATA}}$, $t_{\text{ALL-BANKS-DATA}}$ }
banks-dir	{}	{ $t_{\text{ALL-BANKS-DATA}}$ }
Bank of America-info	{ $t_{\text{BANK-OF-AMERICA}}$ }	{ $t_{\text{BANK-OF-AMERICA}}$ }
US Bank-info	{ $t_{\text{US-BANK}}$ }	{ $t_{\text{US-BANK}}$ }
users-dir	{ $t_{\text{ALL-USERS-DATA}}$ }	{ $t_{\text{ALL-USERS-DATA}}$ }
Bob	{ t_{BOB} }	{ t_{BOB} }
bob-info	{ t_{BOB} }	{ t_{BOB} }
bob-password	{ t_{AUTH} }	{ t_{BOB} }
bob-bankofamerica-creds	{ t_{BOB} , $t_{\text{BANK-OF-AMERICA}}$ }	{ t_{BOB} }
bob-usbank-creds	{ t_{BOB} , $t_{\text{US-BANK}}$ }	{ t_{BOB} }
Dan	{ t_{DAN} }	{ t_{DAN} }
dan-info	{ t_{DAN} }	{ t_{DAN} }
dan-password	{ t_{AUTH} }	{ t_{DAN} }
dan-usbank-creds	{ t_{DAN} , $t_{\text{US-BANK}}$ }	{ t_{DAN} }

Mint File System Hierarchy

Figure 3-2: In continuation of the example in figure 3-1, Bob added accounts in both banks to his Mint account, and Dan added a US BANK account to his Mint account. Tags are shown under the secrecy label and integrity label columns, and are denoted $t_{\text{<name>}}$. t_{BOB} and t_{DAN} are both *USER-DATA* tags.

The application uses Aeolus files for persistent storage. Figure 3-2 shows the hierarchy of the Mint files and their labels, described below:

mint-dir

This is the root directory for the Mint application. The secrecy label is empty, the integrity label contains the *ALL-USERS-DATA*, *ALL-BANKS-DATA* and *ALL-USERS-AUTH* tags.

banks-dir

Contains a *bank-info* file for each bank that was registered to the application. This directory has an empty secrecy label, and has the *ALL-BANKS-DATA* tag in its integrity labels. Each of the *bank-info* files in this directory has a *BANK-DATA* tag in both its secrecy and integrity labels.

users-dir

Has the *ALL-USER-DATA* tag in both its secrecy and integrity labels. Contains a subdirectory per user, each subdirectory contains a *USER-DATA* tag in its secrecy label, and a *USER-DATA* tag in its integrity label. Each of these subdirectories contains the following files:

user-info

Contains the user tag for that user. This file has the *USER-DATA* tag in both its secrecy and integrity labels.

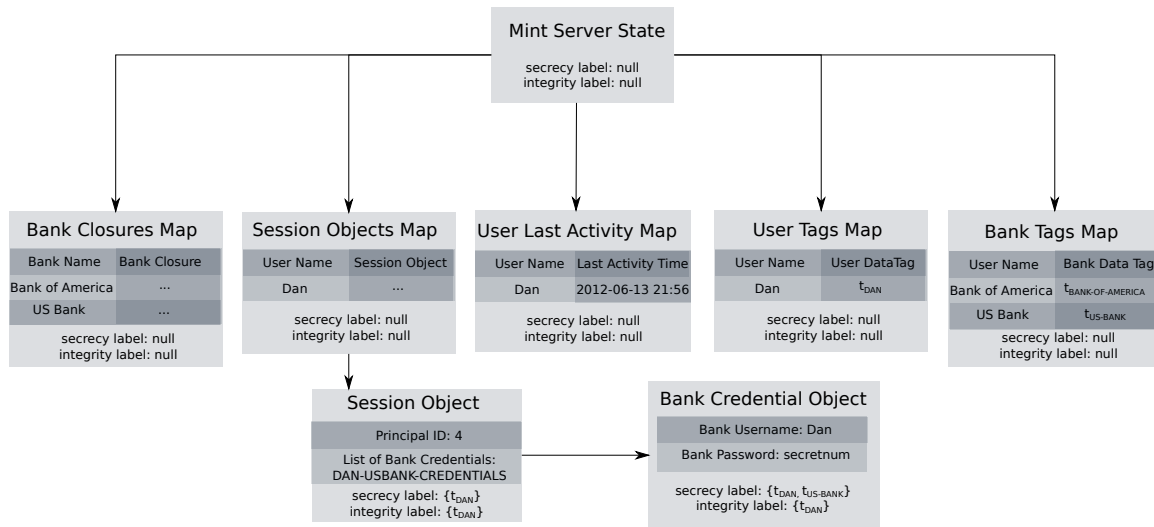
user-password

Contains the user's password. This file contains the *AUTH* tag in its secrecy label, and the *USER-DATA* tag in its integrity label.

user-bank-creds

Contains user credentials for a particular bank. The *users-dir* directory contains one *user-bank-creds* file per bank the user added to their account. This file has the *BANK-DATA* and *USER-DATA* tags in its secrecy label, and the *USER-DATA* tag in its integrity label.

3.1.3 Shared Memory Objects



Mint Shared State Objects

Figure 3-3: This diagram builds on the example presented in Figure 3-2. Bob's session had timed out, and no session object or last activity time is stored for his username.

The application uses Aeolus shared memory objects for storing session state for users and bank closures.

To allow for quick responsiveness to user requests, the application stores some key information in shared memory. Figure 3-3 diagrams the different shared memory object, their relations what information they store.

A *Mint Server State* object that has null labels is stored as the Aeolus *root* object and contains the following objects:

Bank Closures Map

A bank closure is used to retrieve user information from a bank. The bank closure is authoritative for the tag of the bank it is working for. A bank closure has null labels. This object maps each bank name to a bank closure, and has null labels.

Session Objects Map

This object maps a username to a session object. A *session object* contains ses-

sion information about a user. It also contains a list of *bank credential* objects, one for each bank a user has registered for, as well as the user's Principal ID (PID). A *session object* has the *USER-DATA* tag in both its labels. A *bank credential* object contains the user's username and password for that particular bank. It has a *BANK-DATA* tag and a *USER-DATA* tag in its secrecy label, and a *USER-DATA* tag in its integrity label.

User Tags Map

The application stores the tag for each user that has an active session, this allows the application server to access session objects with the user tag in their label. User tags are stored as a mapping from username to tag. This object has null labels.

Bank Tags Map

Bank tags are important to store in memory to allow for quick server response time. They are stored as a map from bank name to bank tag. This object has null labels.

User Last Activity Map

The time of a user's last activity is used to identify which users are still active. This information is stored as a mapping from username to last activity time. This object has null labels.

3.2 Implementation

The application server runs on a VN, and receives user requests via RPC. Users can sign up, login, logout, register a bank and retrieve statistics. A typical flow of RPC calls is shown in figure d.

At startup, the system registers a number of banks. These are the banks which the users can later add accounts for on their Mint account. The system is now ready to accept users.

Users can register to the Mint service through a *signUp* RPC call:

signUp(username, password)

Signs up a new user with the username as *username* and password as *password*. If *username* is available, associates those credentials with a newly-created user data tag and a newly-created principal ID, then stores all this information on disk under a new user directory, as per the file system hierarchy described in section 3.1.2. The system also stores the newly created tag in the *User Tags Map*. Throws an exception if this *username* is already used.

Users can then log in to modify their account and retrieve their information using the following RPC call:

login(username, password)

Authenticates a user, and if successful, stores their information in a session object, and updates their last activity time using the shared memory objects described in section 3.1.3. The system authenticates the credentials by calling an *authentication closure*, which adds the All User Data tag, verifies the credentials, declassifies and terminates if the *username* and *password* are correct, throws an exception otherwise.

Once a user logs in, they only have to include their *username* in persuing requests to carry out different actions. In each of the following RPC calls, authentication is carried out by checking if the provided *username* has a session that has not timed out yet (there is a last activity time associated with a session, and thus must not be older than a certain amount of time, e.g. 15 minutes).

addBank(username, bankName, bankUsername, bankPassword)

Registers the specified bank for the specified username with the given credentials. This request writes a *user-bank-creds* file to disk, with the tag of the user in its integrity label, and the tag of the user and that of the bank in its secrecy label.

downloadTransactions(username)

This action connects to each bank the user has registered for their account,

downloads the latest transactions for the user, processes them, and returns the result. The RPC thread uses the closure of the bank to connect to the bank and declassify the information returned. The thread then adds a user tag to its secrecy label, and uses a reduced authority call to the public principal to process the information; this ensures that information cannot be leaked while it is being processed ².

logout(username)

Terminates the user's session, removing their information from shared memory.

removeBank(username, bankName)

Removes the specified bank from a user's account.

For each RPC call, the thread handling the call starts running on behalf of P_{MINT} , and then reduces its authority as soon as possible to avoid errors leading to information leakage. For example, in *downloadTransactions*, after authenticating the user, the thread reduces its authority to run on behalf of the users principal, then processes the downloaded information on behalf of P_{PUBLIC} and then returns.

3.3 System Security

The use of a *username* in the implementation of the RPC calls defined above to identify which user invoked those calls creates a vulnerability in our system. Alice could send a *downloadTransactions("Bob")* request, and because the information is being sent outside the Aeolus system, the information in the reply would have null labels, and hence Alice could access Bob's information.

Another problem that comes up is related due to the assumption that the user is connecting from another Aeolus node. If this was not the case, i.e. the requests are coming from outside the Aeolus system, sensitive information would be exposed

²One can imagine that the information is being processed by a different application in the system, using this technique means our application does not have to trust that application.

over the network (e.g. the user's Mint credentials), which wouldn't happen for RPCs within the Aeolus system because those are encrypted³.

Both issues can be resolved by using standard cryptography mechanisms to encrypt and authenticate external communication. For example, encrypting all requests and replies coming in and out of the system and, in addition, sessions can be identified by unguessable temporary session-ids, and, those session-ids would be used to authenticate future requests in that session.

³Requests coming from outside the system could be RPC or HTTP requests.

Chapter 4

Auditing the Mint Application

In this chapter we discuss auditing and how the audit trail can be used for discovery.

As described in section 2, every security-related event is logged automatically by Aeolus; in addition, applications can log events of their own, and take advantage of this in the Mint system by logging application-specific events; an example is given in section 4.1.

Events are stored in a table that we will refer to as the *LOG*. The LOG is useful for discovery, which can be done by means of a query over the table. Table 4.1 and figure 4-1 show two examples.

Table 4.1 shows which principals declassified Bob's tag; thus, it shows which user's accessed Bob's information. In a secure implementation of the Mint application, no user principal other than Bob's principal should be able to declassify Bob's tag. However we introduced a bug in our implementation of the Mint application to demonstrate how the LOG can be used for discovery, and hence we can see that many principals, running on behalf of other users, were able to declassify Bob's tag - something that should never happen. In this example, the LOG helps us discover leakage of Bob's information due to an application error.

Figure 4-1 shows a graph of how many user requests caused a declassification of a user data tag for every 5-second period since the application launched; in other words, it shows the activity trend of the user. Any spike in the graph means that

the user was abnormally active in certain periods of time. In this example, the LOG helps us discover if any user accounts were compromised.

secrecy	integrity	user	op_name	timestamp
{11}	{}	Alex	DECLASSIFY	2012-06-13 21:57:06.876
{11}	{}	Aline	DECLASSIFY	2012-06-13 21:57:04.852
{11}	{}	Jack	DECLASSIFY	2012-06-13 21:56:45.73
{11}	{}	Jack	DECLASSIFY	2012-06-13 21:56:42.691
{11}	{}	Mike	DECLASSIFY	2012-06-13 21:56:40.649
{11}	{}	Mark	DECLASSIFY	2012-06-13 21:56:39.624

Bob's Information Leaks

Table 4.1: This table shows the declassifies of Bob's data tag from users other than Bob himself. The columns are as described in section 2.3.5.

In the following section we describe how the data for Figure 4-1 can be extracted from the LOG, as well as the possible uses of summarization for this scenario.

Users Account Activity

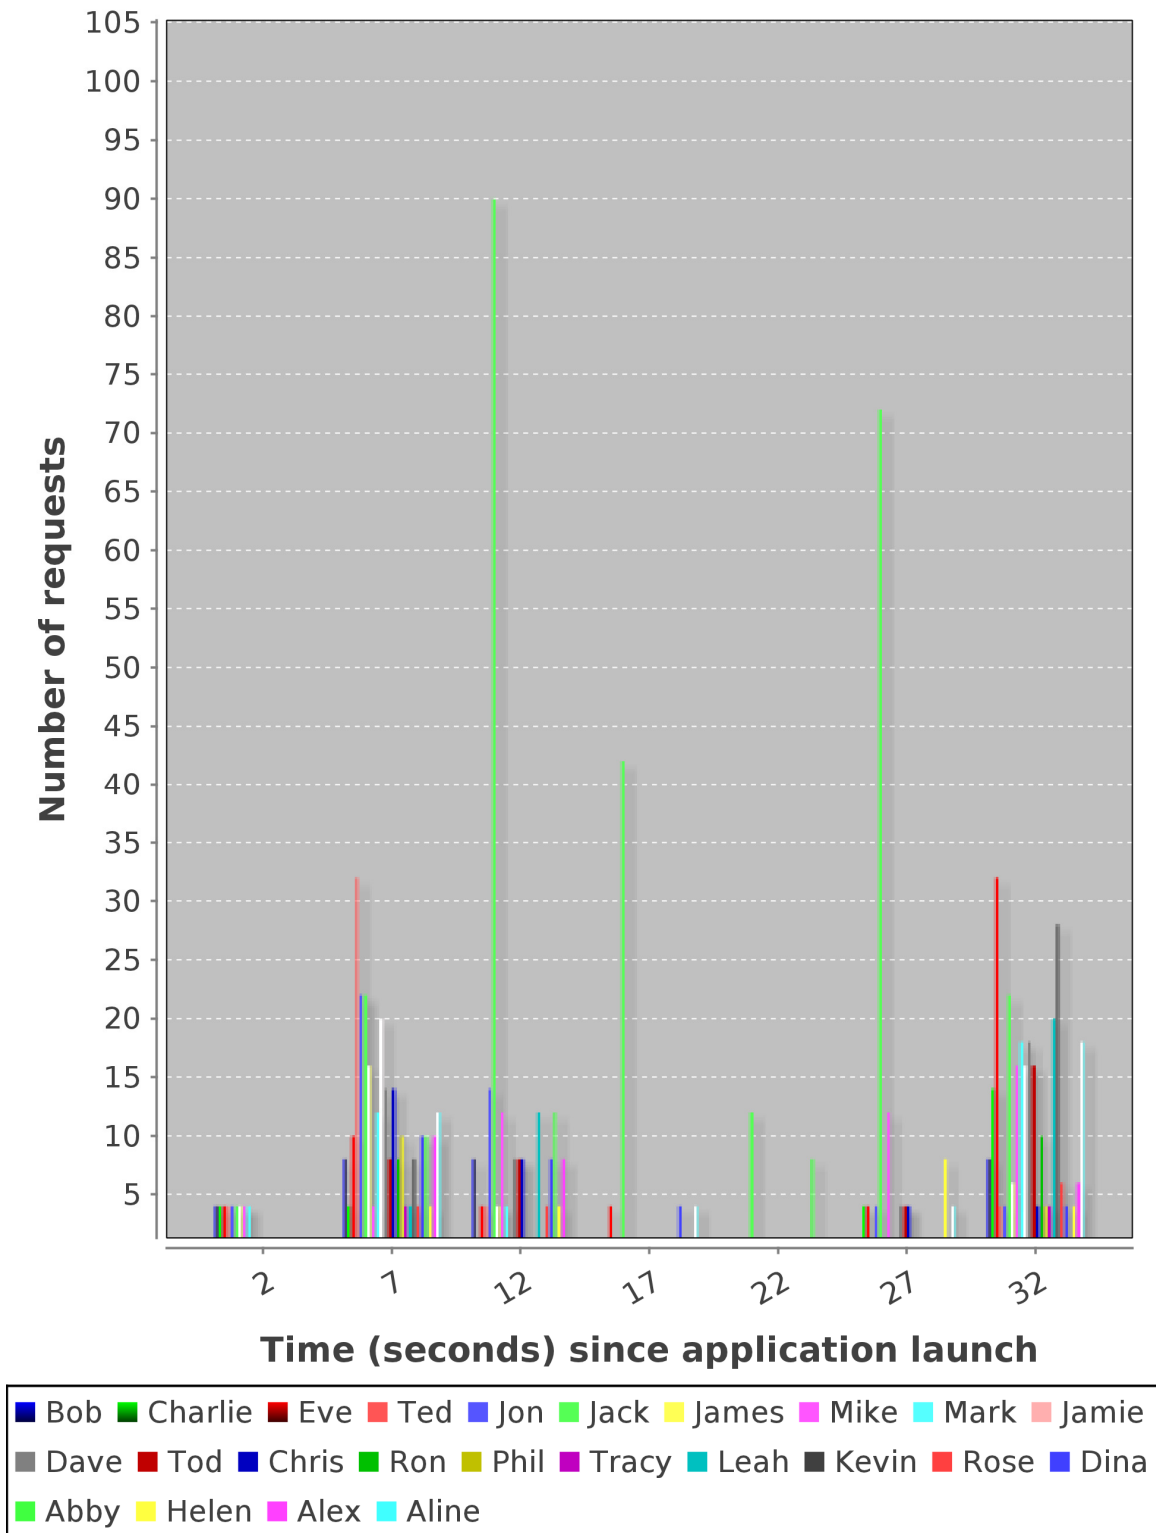


Figure 4-1: This graph shows the activity trend of each user. Notice that the activity for Jack spikes in the graph at multiple places. This could be the result of an attacker accessing his account and issuing requests to quickly retrieve all of Jack's financial information.

4.1 Detecting Suspicious User Activity

In order to gather the information necessary to produce the graph in figure 4-1, we add application-specific events to the LOG for each different action the user takes.

For example, when Bob signs up successfully with our application, we record his username, “Bob”, as well as his data tag and the principal authoritative for that tag. This is done using the following call¹:

```
AeolusLib.createEvent('`mint-user-signup`',  
    Arrays.asList(new String[] {`Bob`, `5`, `10`}));
```

To obtain a table detailing the username, principal, data tag and other sign-up information of all users, we use the following SQL query based on our application-specific events:²:

```
select to_array(app_arg)[0] as pid, to_array(app_arg)[1]  
    as username, to_array(app_arg)[2] as tag, timestamp as  
    signed_up, event_counter from events where  
    op_name='mint-user-signup'
```

Table 4.2 shows a possible result of running the query, with some extra information about the time when each user signed up.

Finally, we can obtain the information necessary to create the graph in figure 4-1 using the following query:

```
select e.count, users.username, e.to_timestamp as  
    timestamp from (select count(*), tags_modified,  
    (to_timestamp(((extract (epoch from  
    events.timestamp)/5)::int)*5)) from events  
    tags_modified in (select tag from users) group by  
    (to_timestamp(((extract (epoch from  
    events.timestamp)/5)::int)*5)), tags_modified) as e  
    inner join users on e.tags_modified=users.tag
```

¹Recording the principal ID and the number of the tag is enough to retrieve the full information about a principal and a tag from the AS.

²`to_array` is a procedure that turns a comma-delimited string into a PSQL array.

pid	username	substring	signed_up	event_counter
6	wissam	4	2012-06-13 21:56:36.207	131
9	david	7	2012-06-13 21:56:36.207	445
12	dan	10	2012-06-13 21:56:37.338	793
13	Bob	11	2012-06-13 21:56:37.338	1148
14	Charlie	12	2012-06-13 21:56:37.338	1211
15	Eve	13	2012-06-13 21:56:37.338	1274
16	Ted	14	2012-06-13 21:56:37.338	1337
17	Jon	15	2012-06-13 21:56:37.338	1400
18	Jack	16	2012-06-13 21:56:37.338	1463
19	James	17	2012-06-13 21:56:37.338	1526
20	Mike	18	2012-06-13 21:56:37.338	1589
21	Mark	19	2012-06-13 21:56:37.338	1652
22	Jamie	20	2012-06-13 21:56:38.578	1715
23	Dave	21	2012-06-13 21:56:38.578	1778
24	Tod	22	2012-06-13 21:56:38.578	1841
25	Chris	23	2012-06-13 21:56:38.578	1904
26	Ron	24	2012-06-13 21:56:38.578	1967
27	Phil	25	2012-06-13 21:56:38.578	2030
28	Tracy	26	2012-06-13 21:56:38.578	2093
29	Leah	27	2012-06-13 21:56:38.578	2156
30	Kevin	28	2012-06-13 21:56:38.578	2219
31	Rose	29	2012-06-13 21:56:38.578	2282
32	Dina	30	2012-06-13 21:56:38.578	2345
33	Abby	31	2012-06-13 21:56:38.578	2408
34	Helen	32	2012-06-13 21:56:38.578	2471
35	Alex	33	2012-06-13 21:56:38.578	2534
36	Aline	34	2012-06-13 21:56:38.578	2597

Users Table

Table 4.2: This is a table detailing the information about users signing up during a simulated run of the Mint application. The *pid* column denotes the user's principal, the *username* column denotes the user's username, the *tag* column denotes the user's data tag, the *signed_up* column denotes the time at which the sign up event took place, and the *event_counter* column denotes the unique event_counter number of the sign up event.

count	username	timestamp
22	Jack	2012-06-13 21:56:40-04
4	Jack	2012-06-13 21:56:35-04
42	Jack	2012-06-13 21:56:50-04
72	Jack	2012-06-13 21:57:00-04
90	Jack	2012-06-13 21:56:45-04
12	Jack	2012-06-13 21:56:55-04
22	Jack	2012-06-13 21:57:05-04

Account activity for user *Jack*

Table 4.3: This table shows the number of declassifies done by Jack’s principal in each five-second period since the application launched. The *count* denotes such number. The *username* column denotes the user who’s data tag is in question. The *timestamp* column denotes the start of the 5-second period. A table containing similar information for all users would be a precursor to producing the graph in figure 4-1, here we only show this data for user Jack for simplicity.

Table 4.3 shows a the result of a modified version of this query that shows information just for the user Jack. This query provides us with a count of how many requests by Jack caused a declassification of a user data tag for every 5-second period of time since the application launch.

4.2 Summarization

The problem with using the LOG for detection is that it gets big. The next chapter addresses this problem.

Chapter 5

The Summary System

This thesis is concerned with providing mechanisms that allow the size of the log to be controlled. A log for a large system (e.g., the Mint application with millions of users) is going to be very big, and the problem gets worse if the system is long-lived. A way to reduce the log is to delete or archive events. Typically, this will be done for events in the past, e.g., events collected last year are less important than recent events. However, there is no way that Aeolus can decide by itself what events are important: this has to be done by the application. A further point is that the application may need to retain some of the information in the events, at least those that are being deleted; again the application has to specify what information should be retained.

For example, if we store the number of declassifies a user's principal carries out per user session, we would require many fewer events to be stored in the log table to formulate a pattern of the user's activity, and, eventually detect suspicious behavior patterns. Thus, we could choose to delete events that have been summarized, without loss of the summarized information.

Our system is designed to support this process. We provide an interface that can be used by application code to create summaries and to mark events for deletion or archiving; from this point forward we refer to these as the *marked* events. Summaries are stored in a *SUMMARY* table and therefore queries can be used to perform discovery on them if desired.

5.1 Summarization Model

An application uses our summary system from within a VN. This has the following ramifications:

- Every access to the LOG or SUMMARY table must obey our information flow constraints. The thread can only read events from these tables whose secrecy labels are contained in those of the thread, and vice versa for integrity labels. Also, all events that are added to the SUMMARY table will have the labels of the thread at the moment the event is added.

The thread producing the summaries or deletes runs with the authority of some principal. For example, a thread creating a summary in the Mint application would likely run the P_{MINT} principal. Thus the thread can use its authority to declassify as needed to produce a summary with the appropriate labels.

- The user code can run an arbitrary Java computation to produce the summary from the events being summarized. This computation could be a simple one, such as a count of the number of the DECLASSIFY events encapsulated by a summary, or a more complex one, such as a categorization of the users activity as normal or suspicious based on the different events encapsulated by the summary.
- The activities of the user code will be audited in the usual way. Thus it will be possible to run discovery over these events, and if desired the user code can log additional events in the base log by using the Aeolus interface.

In the following section, we describe the typical use cases that our system aims to address.

5.2 Summarization Workflow

A typical workflow of how an application might want to use our system to produce a summary is as follows:

1. Possibly run some queries to establish a context for the events to be summarized or marked. For example, identifying Jack's principal ID is necessary to summarize the number of declassifications carried out Jack's principal in a particular session. Such a context could be established by retrieveing a *Users Table* as shown in section 4.1¹.
2. Run a query to identify the events to be summarized or marked. For example, all the declassifications carried out by Jack's principal per session.
3. Compute the information that goes into the summary. For example, whether or not the number of declassifications caused by Jack's principal for a particular session is too high.
4. Possibly produce and store the summary.
5. Possibly mark all events encapsulated by the summary.

The next two sections describe how our system allows applications to query, summarize and mark events.

5.3 The QUERY Table

The first step in producing a summary or marking an event is to identify the events to be summarized or marked. This can be done by running a query. Since these activities are likely to be performed periodically, it is useful to allow the queries to be defined once and then used over and over.

¹Other necessary information for this summary example is to identify the start and end of Jack's sessions.

For example, the query 4.1 could be run many times over the lifetime of the application.

However, an application might need to use a particular query with different parameters - for example, the Mint application might want to run a query to find out the last login time for different users.

A further point is that queries could be run over the LOG or the SUMMARY table, or both². It is up to the application to specifically identify which tables the query will run over.

The following two sections describe the interface we provide to support these use cases, as well as the *QUERY* table which we use to store query records.

5.3.1 The Query Management Interface

We expose the following interface for managing and running queries in the *QUERY* table:

addQuery(name, query_text)

Adds a query with a NAME as *name* and TEXT as *query_text*. If a query with that name already exists, throws an exception.

runQuery(name, vals)

Runs the query with the name *name*, replacing any existing parameter placeholders in the query TEXT with the values in the *vals* array. Finally, the running the query could produce a *view*. This helps in establishing a context for summaries.

Our system logs uses of these functions into the LOG table, detailing function name as the *app_op* and the query name and version as the *app_args*.

We do not support deletion of queries from the *QUERY* table because summaries could refer to them, as described in section 5.4.

A final point is that *runQuery* will return only records that the calling thread is allowed to see per the rules in section 2.2.2.

²Or over user-defined views

5.3.2 Query Attributes

For each record stored in the QUERY table, we store the following attributes:

NAME

This field is provided by the application. It can be used to identify the purpose of the query, for example, query 4.1 could be named “Mint Sign-ups”.

TEXT

This is the text of the query. This does not have to be a complete query: we allow the application to leave query parameters unspecified, replaced with placeholders in this field, so that a record can be re-used with different parameters to represent different queries. The following is an example of what this attribute might hold as a value:

```
select * from LOG where event_counter=?
```

We do not store labels in the QUERY table because we expect those queries to run over the LOG and SUMMARY tables, which already have labels. To avoid any confusion and prevent the use of our interface as a covert channel, we require that a thread have null labels while adding a query record.

5.4 Producing Summaries and Marking Events

An application can use queries in the QUERY table to set up a context for summarization and identify events to be summarized. The application can then store a summary of those events in the SUMMARY table. The application might then choose to mark the events encapsulated by a particular summary event, or summarize a group of summary events.

For example, a summary of the information presented in figure 4.1 would detail all the users other than Bob that accessed Bob’s information, along with the period of time during which those accesses happened (identified by a start event_counter and an end event_counter).

The following two sections describe the interface we provide to support producing summaries and marking events, as well as the SUMMARY table which we use to store summary records.

5.4.1 SummaryObjects

In this section we describe *SummaryObjects*. *SummaryObjects* provide an interface for applications to create and add summaries to the SUMMARY table, as well as mark their events.

The interface we expose is as follows:

SummaryObject(query_name, query_params)

Retrieves the events identified by the query *query_name* with parameters *query_params*.

getEvents()

Returns the events encapsulated by this SummaryObject.

addSummary(summary_name, summary_args)

Adds the summary represented by this *SummaryObject* to the SUMMARY table with name as *summary_name* and summary ARGS as *summary_args*. Stores the calling threads' secrecy and integrity labels as the summary's *SECRECY* and *INTEGRITY*.

mark(constraint_query, query_params)

Marks the group of events underlying this *SummaryObject*. We prevent certain Aeolus events from being marked. We also allow applications to specify which events should not be marked for a particular call to this function, using the *constraint_query* and *query_params* arguments.

In our scheme, we depend on being able to identify events underlying a particular summary by their *event_counter*. Hence we require all queries used to produce *SummaryObjects* to include an *event_counter* attribute.

Each use of a *SummaryObject* method is logged to the LOG table. We log the name of the function along with the arguments to identify the exact actions of the application.

5.4.2 Summary Attributes

For each record stored in the SUMMARY table, we store the following attributes:

EVENT_COUNTER

Unique identifier for each record.

NAME

Similar to *NAME* in the QUERY table, this field is provided by the application and could be used to identify the content of the summary record.

SECRECY

This is the secrecy label of the thread that added the summary record.

INTEGRITY

This is the integrity label of the thread that added the summary record.

QUERY_NAME

This field is provided by the application. It identifies the query in the QUERY table that identifies the events summarized by this record.

QUERY_PARAMS

This field is provided by the application. It stores the parameter values for the query identified by the *QUERY_ID* if any.

ARGS

This field is provided by the application. It stores the application's computation of the summary. In the example mentioned above, this could be the list of users accessing Bob's information over the period of time elapsed between the event identified by *START_COUNTER* and the event identified by *END_COUNTER*.

TIMESTAMP

The real time at which this summary was created.

The application only needs to specify a query name and parameters to identify the events to be summarized, and a name for the summary in order to add a summary to the SUMMARY table. The rest of the attributes are computed internally by our system.

5.5 Marking Events

Our design allows but does not require the production of a summary event in conjunction with marking. For instance, a *SummaryObject* can be created, then used to mark events, without the application having to call *addSummary*.

Since our system was not intended to alter the LOG or SUMMARY table, we keep track of marked events in two additional tables: LOG-MARKED and SUMMARY-MARKED. Those tables include events that have been marked by the application.

One can imagine that our system runs in a *hide-mode*. It hides events that have been marked from the application. Those events are not deleted, as mentioned before, and can could be made visible to the user in future work [?].

Finally, we reiterate that our system only provides a mechanism to identify marked events, and leaves the question of how to handle those events up to the application.

5.6 Discussion

It is possible through our system that an application summarizes pre-existing summaries. For example, the Mint application, after summarizing declassify events for user data tags per user session, could use those summaries to summarize declassify events for user data tags per month.

Our system supports this functionality since it does not make any strong assumption about queries in the QUERY table. This allows the application to treat events in the LOG or SUMMARY table the same way.

Finally, by allowing applications to protect events from being marked, and internally preventing important Aeolus events from being marked, our system provides a way to prevent application errors marking important information.

Chapter 6

Auditing Using Summarization

In this chapter we visit examples of how our system can be used to summarize and access information for the Mint application, as well as mark certain events.

6.1 Summarizing User Sessions

In this example we will show how to summarize user actions per session. For example, for a particular session for Bob, we will record which of the user actions presented in section 3.2 did he call.

This should not be a hard task since we use application-level logging to track all of users' actions. The trick, however, is to identify user sessions encapsulating these actions.

First, we initialize the summarization system, which we will use to add and run queries:

```
//SummarySystem uses the singleton paradigm  
SummarySystem system = SummarySystem.getInstance();
```

We then create a view that contains relates users to their principal ID and data tag, as shown in 4.2:

```
String users_query = ``create or replace view users as  
    select running_principal as pid, app_name as username,
```

```

app_arg as tag, timestamp as signed_up, event_counter
from events where op_name=1000 and app_name <>'login'
and app_name <> 'logout' and app_name <> 'userAttack'
and app_name <> 'collectReport' and app_name <>
'addBank''';

```

```

system.addQuery('`users-info`', users_query);
system.runQuery('`users-info`', null);

```

To do so, we first identify user the times at which users logged in and out:

```

// When a user, say Bob, logs in, we execute the
// following code:
// AeolusLib.createEvent('`mint-user-login`',
// Arrays.asList(new String[]{`Bob`}));
// and similarly for logouts.

```

```

String logins_query = ``create or replace view logins as
select app_arg[0] as username, timestamp as start,
event_counter from events where
app_op='mint-user-login''';

```

```

String logouts_query = ``create or replace view logouts as
select app_arg[0] as username, timestamp as start,
event_counter from events where
app_op='mint-user-login''';

```

```

//add the queries
system.addQuery('`logins`', logins_query);
system.addQuery('`logouts`', logouts_query);

```

```
//creates the login and logout views
system.runQuery('`logins`', null);
system.runQuery('`logouts`', null);
```

Now we can use this information to identify user sessions:

```
String sessions_query = `select logins.username,
    logins.event_counter as start_counter,
    min(logouts.event_counter) as end_counter,
    logins.start, min(logouts.end) as end from logins inner
    join logouts on logins.username=logouts.username where
    logouts.event_counter>logins.event_counter group by
    logins.username, logins.event_counter, logins.start`;

system.addQuery('`sessions`', sessions_query);
```

Now we can use this query to retrieve the start and end times of all user se

In order to summarize which actions a user carried out during each of their sessions, we execute the following code:

```
String user_session_actions_query = `select app_op from
    events where app_op like 'mint-user-\%' and app_arg[0]
    = ? and event_counter > ? and event_counter < ?`;

system.addQuery('`user-sessions-actions`',
    user_sessions_actions_query);
```

The *user-sessions-actions* query will provide the different operations carried out by a user, the user name, and the session's start counter for all users and all sessions. Now we would like to add a summary for each session that contains the username as well as the actions carried out:

```
ResultSet rs = system.runQuery('`sessions`', null);

while(rs.next){
```

username	start_counter	end_counter
wissam	168	373
david	482	1075
dan	830	1086
Ted	2644	22675
Leah	2690	19437
Alex	2736	7786
Jamie	2782	5991
Dina	2828	7336
Mark	2874	18576
Helen	2920	6312
Jon	2966	18370
Mike	3072	18321
Jack	3118	15342
Eve	3164	7735
Charlie	3210	7434
Phil	3450	19535
Tod	3684	5895
Chris	3924	18075
James	3970	7635
Dave	4110	18223
Aline	4253	18773
Rose	4580	17816
Ron	4626	19699
Kevin	4891	17865
Tod	5931	17714

User Sessions Table

Table 6.1: This table shows start and end event counters for different user sessions. Start and end times are omitted for simplicity.

```

String username = rs.getString(1);
long start_counter = rs.getLong(2);
long end_counter = rs.getLong(3);
Object[] vals = new Object[] {username, start_counter,
    end_counter};
SummaryObject so = new
    SummaryObject(``user-session-actions``, vals);

//process is a function that computes a summary
//for the events identified by a query. For now, we can
assume
//it aggregates as a string all the user actions
identified.
String summary = process(so.getEvents());
so.addSummary(``user-session-summary``,
    String.format(``\%s, \%d, \%d``, username,
        start_counter, end_counter, summary));
}

```

Finally, we summarized user actions per session by adding records to the SUMMARY table which contain information about the user, start and end counters of the session, as well as all the user actions that took place in that session.

6.2 Summarizing File System Events

In this example we explore how to summarize information about file system events; we show how to store a list of added/removed and modified files in a week leading up to a certain time.

```

SummarySystem system = SummarySystem.getInstance();

String fs_access = ``select * from events where
    filename<>'' and DATE_SUB(?, INTERVAL 7 DAY)``;

```

```

system.addQuery('`fs_week_access`', fs_access);

long now = System.currentTimeMillis();
SummaryObject so = new SummaryObject('`fs_week_access`',
    Object[] {new java.sql.Timestamp(now)});

String summary = process(so.getEvents());
so.addSummary('`File accesses`', String.format('`%d,
    %s`', now, summary));

```

This will store a summary of the file system events in the last week. Only file system events use the *filename* attribute in the LOG, hence we use that to retrieve all file system events.

We might wish to mark most file system events, except for the ones that change the file system hierarchy (adding or removing directories). We could do this as follows:

```

String fs_protect = '`select * from events where op_name =
    ? or op_name = ?`'

system.addQuery('`protect-fs-events`', fs_protect);

so.mark('`protect-fs-events`', new
    Object[] { '`CREATE_DIRECTORY`', '`REMOVE_DIRECTORY`' });

```

6.3 Summarizing User Trends

As a final example, we will show how records in the SUMMARY table itself could be summarized. In this section we will further summarize session summaries, explained in section 6.1, and coalesce them into summaries that

```

SummarySystem system = SummarySystem.getInstance();

```

```

String session_summary_query = ``select * from SUMMARY
    where name=user-session-summary and end < ?`;

// This time tells us when did the application launch and
    when the last
// event took place.
String min_event_time = ``select min(timestamp) from
    events``;
String max_event_time = ``select max(timestamp) from
    events``;
system.addQuery(min_event_time);
system.addQuery(max_event_time);

java.sql.Timestamp min =
    system.runQuery(min_event_time).getTimestamp(1);
java.sql.Timestamp max =
    system.runQuery(max_event_time).getTimestamp(1);

// Now we can run our summary of summaries for every week
    since application
// launch until the last event.

system.addQuery(``protect-fs-events``, fs_protect);

so.mark(``protect-fs-events``, new
    Object[]{'CREATE_DIRECTORY', 'REMOVE_DIRECTORY'});

```

6.4 Sufficiently Summarized Information

The summaries produced in sections 6.1 and 6.2 overlap in what they summarize. For example, adding a *user-bank-creds* file to the file system happens only when a user adds a bank to their account, i.e. call the *addBank(...)* request. Hence this particular piece of information could be derived from two summaries: one that summarizes the user's actions, the other summarizing the file system operations.

There are many other examples in which summaries could overlap in what they summarize, and it is one of the reasons why our system chooses to leave it up to the application to decide when certain events are “sufficiently summarized” and can hence be marked, rather than trying to determine which events have already been summarized and will not be needed anymore.

Bibliography

- [1] Aaron Blankstein. Analyzing audit trails in the Aeolus security platform. Master's thesis, MIT, Cambridge, MA, USA, June 2011.
- [2] Winnie Cheng, Dan R. K. Ports, David Schultz, Victoria Popic, Aaron Blankstein, James Cowling, Dorothy Curtis, Liuba Shrira, and Barbara Liskov. Abstractions for usable information flow control in Aeolus. In *Proceedings of the 2012 USENIX Annual Technical Conference*, Boston, MA, USA, June 2012.
- [3] Winnie Wing-Yee Cheng. *Information Flow for Secure Distributed Applications*. Ph.D., MIT, Cambridge, MA, USA, August 2009. Also as Technical Report MIT-CSAIL-TR-2009-040.
- [4] F. Peter McKee. A file system design for the aeolus security platform. Master's thesis, MIT, Cambridge, MA, USA, September 2011.
- [5] Victoria Popic. Audit trails in the Aeolus distributed security platform. Master's thesis, MIT, Cambridge, MA, USA, September 2010. Also as Technical Report MIT-CSAIL-TR-2010-048.